

SATURN: a library of verified concurrent data structures for OCAML 5

Clément Allain (INRIA)
Vesa Karvonen (Tarides)
Carine Morel (Tarides)

May 30, 2024

1 Abstract

We present SATURN, a new OCAML 5 library available on [opam](#). SATURN offers a collection of concurrent efficient data structures: stack, queue, skiplist, hash table, work-stealing deque, etc. It is well tested, benchmarked and in part formally verified.

2 Motivation

Sharing data between multiple threads or cores is a well-known problem. A naive approach is to take a sequential data structure and protect it with a lock. However, this is often not the best solution. First, if performance is a concern, this approach is likely to be inefficient as locks introduce significant contention. Second, it may not be a sound solution as it can lead to liveness issues such as deadlock, starvation and priority inversion.

In contrast, *lock-free* implementations relying on fine-grained synchronisation rather than locks are typically faster and guaranty system-wide progress. Yet, they are also more complex and come with their own set of bugs: ABA problem (largely mitigated in garbage-collected languages), data races, unexpected behaviors due to non-linearizability.

In this context, SATURN provides a collection of standard lock-free data structures saving OCAML 5 programmers the trouble of designing their own.

3 Library design

SATURN aims at covering a wide range of use cases, from simple stacks and queues to more complex data structures like skiplists and hash tables. More precisely, it currently features: (A) numerous queues: a queue based on the well-known Michael-Scott queue [2], a single-producer single-consumer queue, a multiple-producer single-consumer queue and a bounded queue; (B) a stack based on the Treiber stack [1]; (C) a work-stealing deque; (D) a bag; (E) a hash table; (F) a skiplist.

Most implementations are based on well-known algorithms. They have been adapted to work with and take advantage of the OCAML 5 memory model. For instance, we had to rework the Michael-Scott queue to avoid a memory leak.

Regarding performance, we are working on providing benchmarks for each SATURN's data structure, and significant effort has been dedicated to micro-optimization. In particular, we worked on (a) preventing false sharing¹ (b) adding fenceless atomic reads when possible, which improves performance on ARM processors and (c) avoiding the extra indirection in arrays of atomics to reduce memory consumption.

¹False sharing occurs when different domains access different data items contained in the same cache line, forcing unnecessary synchronization. To prevent this, these data must be padded to ensure they are not in the same cache line.

The feedback we produced while optimicing Saturn has highlighted some missing features in OCAML 5 and led to improvements in upstream OCAML ([padded atomics](#), [CSE bug fixed](#)).

To explore some of these optimizations, we use unsafe features of the language (e.g., `Obj.magic`). Although we design our code to be memory-safe under regular use (e.g. only one domain pushes at any given times in a single consumer single producer queue), we cannot offer the same level of guarantee as with the standard implementations. Consequently, some of [SATURN](#)’s data structures have two versions: (1) a version that does not use any unsafe features of OCAML and (2) an optimized version. While most users should find the regular version efficient enough for their needs, more adventurous users may prefer the optimized version, provided they encapsulate it correctly and thoroughly verify their code.

4 Benchmarks

As we are still in the experimental phase, we only provide some rough preliminary numbers to give an idea of the library’s performance. The following table shows the throughput of three implementations of a single consumer single producer queue in million messages per second. The first implementation is the one provided by [SATURN](#) under the name [Single_cons_single_prod_queue](#), the second one is an optimized version that used unsafe feature of OCAML and the last one is written in [C++ implementation](#) and provide a good upper limit. The tests were run on an Intel i7-1270P and an Arm Cortex M3.

Throuput (million message/s)	Intel	Arm
Spsc queue	60 M/s	100 M/s
Optimized spsc queue	100 M/s	250 M/s
C++ spsc queue	400 M/s	500 M/s

5 Tests

In multicore programming, it is essential to test not only the safety of the data structures but also to verify linearizability and lock-freedom when expected. To achieve this, [SATURN](#) has been thoroughly tested using two primary different tools: [DSCHECK](#) and [STM](#).

[STM](#) is used not only for unit testing but also for linearizability. It automatically generates random full programs using the provided API—in the case of [SATURN](#), a data structure. These programs are executed in parallel with two domains and all results are checked against the post-conditions of each function, providing unit testing. Simultaneously, [STM](#) verifies linearizability by ensuring that all intermediate states can be explained by a sequential execution of the calls. The [STM test for the Treiber stack](#) are a good example of how simple this is to write.

[DSCHECK](#) is a model checker based on the DPOR² algorithm [3]. It is designed to compute all possible interleavings of a given program and verify that each one returns the expected result. This is particularly useful for catching elusive bugs that occur only in specific, rare interleavings. Additionally, [DSCHECK](#) can be used to verify that a program is lock-free, as it will fail to terminate if any form of blocking is present. This is a bit more cumbersome to use than [STM](#) (see the [DSCheck tests for the Treiber stack](#)) but it is still a powerful tool. The implementation has been optimized² to make the tests quick enough to be used even on the more complex data structures of [SATURN](#) (see the [skiplist DSCheck tests](#)).

²DPOR stands for Dynamic Partial-Order Reduction

²See the PRs about [source sets](#) and [granular dependency relation](#).

6 Formal verification

Lock-free algorithms are notoriously difficult to get right. To provide stronger guarantees, we have verified part of **SATURN**'s data structures and aim at covering the entire library. Moreover, one important benefit is that we get formal specifications for verified data structures.

This verification effort has been conducted using **IRIS** [4], a state-of-the-art mechanized *concurrent separation logic*. In particular, all proofs are formalized in **Coq**.

A common criterion to specify concurrent operations on shared data structures is *linearizability*. The equivalent **IRIS** notion is *logical atomicity*: there exists a point in time when a concurrent operation atomically takes effect (the linearization point). This statement takes the form of an *atomic specification*:

$$\frac{\frac{\frac{\{ \text{queue-inv } t \}}{\langle \forall vs. \text{queue-model } t \text{ } vs \rangle}}{\text{queue_push } t \text{ } v}}{\langle \text{queue-model } t \text{ } (vs \# [v]) \rangle}}{\{ () . \text{True} \}}$$

In this example, we specify the `queue_push` operation from an implementation of a concurrent queue. Similarly to **Hoare triples**, the two assertions inside curly brackets express the precondition and postcondition. Here, the `queue-inv t` precondition represents the queue invariant. As it is persistent, we do not need to give it back in the postcondition. The other two assertions inside angle brackets express the *atomic precondition* and *atomic postcondition*. Basically, they specify the linearization point of the operation: during the execution of `queue_push`, the logical model of the queue is atomically updated from `vs` to `vs # [v]`, in other words `v` is atomically pushed at the back of the queue.

As a final note, we emphasize that our verification assumes a sequentially consistent memory model. Nevertheless, OCAML 5's relaxed memory model has been formalized [5] in **IRIS**. It should be possible and is future work to adapt our specifications and proofs to support it.

7 Talk proposal

In our talk, we will introduce **SATURN**, including the main data structures and design guidelines. We will also discuss ongoing work on verifying the library using the **IRIS** concurrent separation logic.

References

- [1] Robert K. Treiber. *Systems programming: Coping with parallelism*. Tech. rep. RJ 5118. IBM Almaden Research Center, 1986.
- [2] Maged M. Michael and Michael L. Scott. “Simple, fast, and practical non-blocking and blocking concurrent queue algorithms”. In: *Proceedings of the fifteenth annual ACM symposium on Principles of distributed computing* (1996), pp. 267–275. URL: https://www.cs.rochester.edu/u/scott/papers/1996_PODC_queues.pdf.
- [3] Cormac Flanagan and Patrice Godefroid. “Dynamic partial-order reduction for model checking software”. In: *Proceedings of the 32nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL '05. Long Beach, California, USA: Association for Computing Machinery, 2005, pp. 110–121. ISBN: 158113830X. URL: <https://doi.org/10.1145/1040305.1040315>.
- [4] Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek Dreyer. “Iris from the ground up: A modular foundation for higher-order concurrent separation logic”. In: *J. Funct. Program.* 28 (2018), e20. URL: <https://doi.org/10.1017/S0956796818000151>.

- [5] Glen Mével, Jacques-Henri Jourdan, and François Pottier. “Cosmo: a concurrent separation logic for multicore OCaml”. In: *Proc. ACM Program. Lang.* 4.ICFP (2020), 96:1–96:29. URL: <https://doi.org/10.1145/3408978>.