

Saturn: a library of verified concurrent data structures for OCaml 5

Clément Allain (INRIA Paris)

Vesa Karvonen (Tarides)

Carine Morel (Tarides)

September 7, 2024

Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
 - ◇ well-tested
 - ◇ benchmarked
 - ◇ optimized
 - ◇ verified

Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
 - ◇ well-tested
 - ◇ benchmarked
 - ◇ optimized
 - ◇ verified
- Writing concurrent code is hard

Buggy queue with size

```
module Queue = Saturn.Queue

type 'a t = {size : int Atomic.t; queue : 'a Queue.t}

let create () =
  {size = Atomic.make 0; queue = Queue.create ()}

let push t msg =
  Atomic.incr t.size;
  Queue.push t.queue msg

let pop_opt t =
  match Queue.pop_opt t.queue with
  | Some elt -> Atomic.decr t.size; Some elt
  | None -> Atomic.set t.size 0; None

let size t = Atomic.get t.size
```

Buggy queue with size

```
let test () =  
  let queue = create () in  
  let d1 = Domain.spawn (fun () -> push queue 1) in  
  let d2 = Domain.spawn (fun () -> pop_opt queue |>  
    ignore) in  
  Domain.join d1;  
  Domain.join d2;  
  pop_opt queue |> ignore;  
  size queue
```

Buggy queue with size

```
let test () =  
  let queue = create () in  
  let d1 = Domain.spawn (fun () -> push queue 1) in  
  let d2 = Domain.spawn (fun () -> pop_opt queue |>  
    ignore) in  
  Domain.join d1;  
  Domain.join d2;  
  pop_opt queue |> ignore;  
  size queue
```

In 10 to 20% of the tries, the test returns a size of -1 .

Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
 - ◇ well-tested
 - ◇ benchmarked
 - ◇ optimized
 - ◇ verified
- Writing concurrent code is hard
 - ◇ bugs that can be hard to reproduce and understand

Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t

let create () = Atomic.make []

let rec push q a =
  let old = Atomic.get q in
  if Atomic.compare_and_set q old (a :: old) then ()
  else push q a

let rec pop_opt q =
  let old = Atomic.get q in
  match old with
  | [] -> None
  | x :: xs ->
    if Atomic.compare_and_set q old xs then Some x
    else pop_opt q
```


Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t
```

```
let cr
```

```
let re
```

```
let
```

```
if A
```

```
else
```

```
let re
```

```
let
```

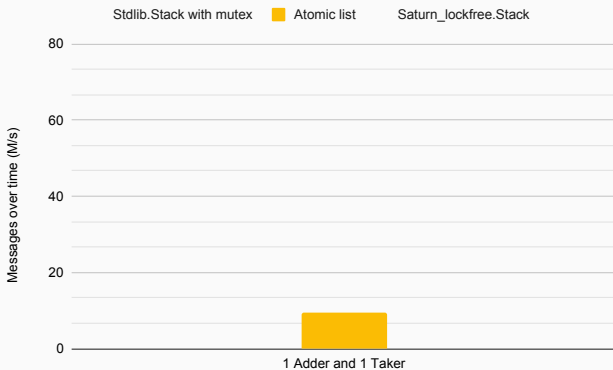
```
match
```

```
| []
```

```
| x
```

```
if Atomic.compare_and_set q old xs then Some x
```

```
else pop_opt q
```



Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t
```

```
let cre
```

```
let rec
```

```
let c
```

```
if At
```

```
else
```

```
let rec
```

```
let c
```

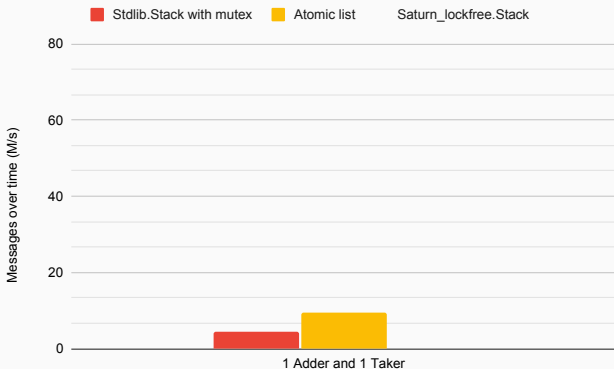
```
match
```

```
| []
```

```
| x :
```

```
if Atomic.compare_and_set q old xs then Some x
```

```
else pop_opt q
```



Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t
```

```
let crea
```

```
let rec
```

```
let oi
```

```
if Atc
```

```
else {
```

```
let rec
```

```
let oi
```

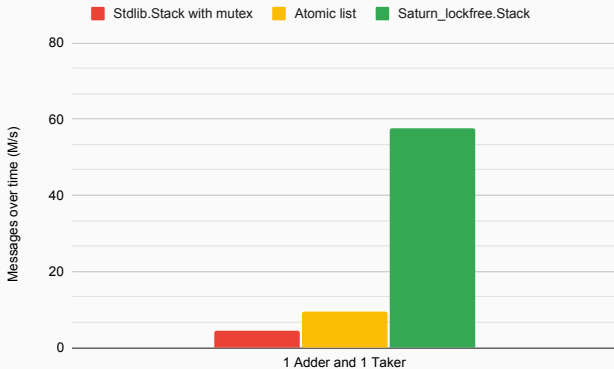
```
match
```

```
| [] -
```

```
| x :
```

```
if Atomic.compare_and_set q old xs then Some x
```

```
else pop_opt q
```



Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
 - ◇ well-tested
 - ◇ benchmarked
 - ◇ optimized
 - ◇ verified
- Writing concurrent code is hard

What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
 - ◊ Queues

What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
 - ◇ Queues
 - multi-producer, multi-consumer
(based on Michael-Scott queue algorithm)
 - single-producer, single-consumer
 - single-producer, multi-consumer
 - bounded, blocking (opened PRs)

What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
 - ◊ Queues

What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
 - ◊ Queues
 - ◊ Work-stealing deque
 - ◊ Stacks
 - ◊ Hashtable (almost in a PR)
 - ◊ Skiplist: a sorted linked list with $o(\log(n))$ operations
 - ◊ Bag

What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
 - ◊ Queues
 - ◊ Work-stealing deque
 - ◊ Stacks
 - ◊ Hashtable (almost in a PR)
 - ◊ Skiplist: a sorted linked list with $o(\log(n))$ operations
 - ◊ Bag
- Need something else ?
 - Open an issue on [ocaml-multicore/saturn](#)

What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
 - ◊ Queues
 - ◊ Work-stealing deque
 - ◊ Stacks
 - ◊ Hashtable (almost in a PR)
 - ◊ Skiplist: a sorted linked list with $o(\log(n))$ operations
 - ◊ Bag
- Need something else ?
 - Open an issue on [ocaml-multicore/saturn](#)
- Most available data structures are lock-free
 - Two libraries: SATURN and SATURN_LOCKFREE

Testing in Saturn

What should (and can) be tested ?

- Correctness
- Linearizability
- Progress (i.e. lock freedom)

How ?

- Lin / STM (ocaml-multicore/multicoretests):
 - ◇ Correctness
 - ◇ Linearizability
- Dscheck (ocaml-multicore/dscheck):
 - ◇ Correctness
 - ◇ Lock-freedom

About optimizations

Numerous micro-optimizations:

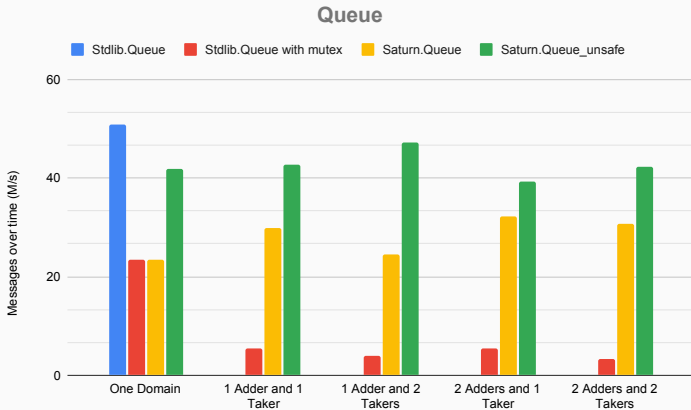
- ◇ Improving algorithms,
- ◇ Preventing false sharing,
- ◇ Removing indirections etc..
- ◇ Experimental: some optimizations use **Obj.magic** (e.g. `Saturn.Queue_unsafe`)

About optimizations

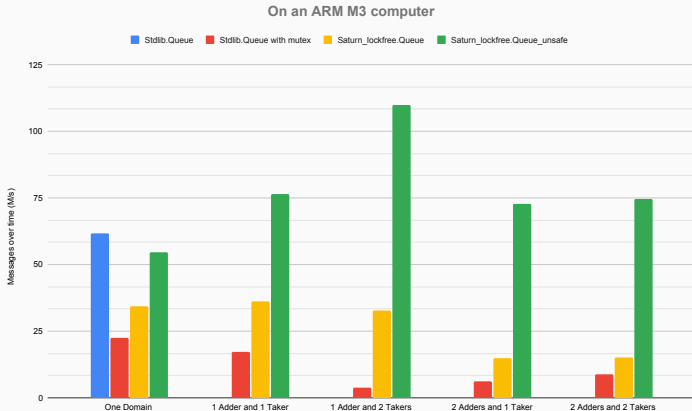
Numerous micro-optimizations:

- ◇ Improving algorithms,
- ◇ Preventing false sharing,
- ◇ Removing indirections etc..
- ◇ Experimental: some optimizations use **Obj.magic** (e.g. `Saturn.Queue_unsafe`)
 - a source of inspiration for language features for efficient concurrent code (e.g. opened PR on atomic record fields)

Benchmarks



Benchmarks



Testing a concurrent algorithm is hard due to the number of potential interleavings.

Formal verification is here to help us!



From OCaml to Coq

```
let rec push t v =  
  let old = Atomic.get t in  
  let new_ = v :: old in  
  if not (Atomic.compare_and_set t old new_) then (  
    Domain.cpu_relax () ;  
    push t v  
  )
```

OCAML

```
Definition stack_push : val :=  
  rec: "stack_push" "t" "v" =>  
    let: "old" := !"t" in  
    let: "new" := `Cons( "v", "old" ) in  
    ifnot: CAS "t" "old" "new" then (  
      Yield ;;  
      "stack_push" "t" "v"  
    ) .
```

Coq

$$\text{linearizability} \quad \underset{\text{in } \overline{\text{SC}}}{\approx} \quad \frac{\frac{\frac{\{ \text{stack-inv } t \}}{\langle \text{vs. stack-model } t \text{ vs } \rangle}}{\text{stack_push } t \text{ } v}}{\langle \text{stack-model } t (v :: \text{vs}) \rangle}}{\{ () . \text{True} \}}$$

Theorems for free from separation logic specifications

Birkedal, Dinsdale-Young, Guéneau, Jaber, Svendsen & Tzevelekos

Writing concurrent protocols in Iris

Definition `stack_inv t ι : iProp Σ :=`
 `\exists l γ ,`
 `\ulcorner t = #l \urcorner * meta l nroot γ *`
 `inv ι (`
 `\exists vs, l \mapsto lst_to_val vs * stack_model2 γ vs`
 `).`

Lemma `stack_push_spec t ι v :`
 `<<< stack_inv t ι`
 `| \forall vs, stack_model t vs >>>`
 `stack_push t v @ $\uparrow\iota$`
 `<<< stack_model t (v :: vs)`
 `| RET (); True >>>.`

Proof.

...

Qed.



Thank you for your attention!

Relaxed memory model (future work)

$$\frac{\frac{\frac{\{ \text{stack-inv } t \}}{\langle v_0, \dots, v_n, \mathcal{V}_0, \dots, \mathcal{V}_n. \text{stack-model } t ((v_0, \mathcal{V}_0), \dots (v_n, \mathcal{V}_n)) \rangle}}{\text{stack_push } t \ v, \mathcal{V}}}{\langle \text{stack-model } t ((v, \mathcal{V}), (v_0, \mathcal{V}_0), \dots (v_n, \mathcal{V}_n)) \rangle}}{\{ () . \text{True} \}}$$

Cosmo: a concurrent separation logic for multicore OCAML,
Mével, Jourdan & Pottier

*Formal verification of a concurrent bounded queue in a weak
memory model,* Mével & Jourdan

Stack benchmarks

