

# Saturn: a library of verified concurrent data structures for OCaml 5

---

Clément Allain (INRIA Paris)

Vesa Karvonen (Tarides)

Carine Morel (Tarides)

September 6, 2024

- ☐ Why SATURN?
- ☐ What is in SATURN?
- ☐ Optimizations and benchmarks
- ☐ Testing
- ☐ Verification

# Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
  - ◇ well-tested
  - ◇ benchmarked
  - ◇ optimized
  - ◇ verified

# Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
  - ◇ well-tested
  - ◇ benchmarked
  - ◇ optimized
  - ◇ verified
- Writting concurrent code is hard

## Buggy queue with size

```
module Queue = Saturn.Queue

type 'a t = {size : int Atomic.t; queue : 'a Queue.t}

let create () =
    {size = Atomic.make 0; queue = Queue.create ()}

let push t msg =
    Atomic.incr t.size;
    Queue.push t.queue msg

let pop_opt t =
    match Queue.pop_opt t.queue with
    | Some elt -> Atomic.decr t.size; Some elt
    | None -> Atomic.set t.size 0; None

let size t = Atomic.get t.size
```

## Buggy queue with size

```
let test () =  
  let queue = create () in  
  let d1 = Domain.spawn (fun () -> push queue 1) in  
  let d2 = Domain.spawn (fun () -> pop_opt queue |>  
    ignore) in  
  Domain.join d1;  
  Domain.join d2;  
  pop_opt queue |> ignore;  
  size queue
```

## Buggy queue with size

```
let test () =  
  let queue = create () in  
  let d1 = Domain.spawn (fun () -> push queue 1) in  
  let d2 = Domain.spawn (fun () -> pop_opt queue |>  
    ignore) in  
  Domain.join d1;  
  Domain.join d2;  
  pop_opt queue |> ignore;  
  size queue
```

In 10 to 20% of the tries, the test returns a size of  $-1$ .

# Why Saturn?

Github: [ocaml-multicore/saturn](https://github.com/ocaml-multicore/saturn)

- A collection of concurrent-safe data structures for OCAML 5:
  - ◇ well-tested
  - ◇ benchmarked
  - ◇ optimized
  - ◇ verified
- Writing concurrent code is hard
  - ◇ bugs that can be hard to reproduce and understand



## Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t

let create () = Atomic.make []

let rec push q a =
  let old = Atomic.get q in
  if Atomic.compare_and_set q old (a :: old) then ()
  else push q a

let rec pop_opt q =
  let old = Atomic.get q in
  match old with
  | [] -> None
  | x :: xs ->
    if Atomic.compare_and_set q old xs then Some x
    else pop_opt q
```

# Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t
```

```
let create
```

```
let rec push
```

```
  let old =
```

```
  if Atomic
```

```
  else push
```

```
let rec pop
```

```
  let old =
```


```
  match old
```

```
  | [] -> None
```

```
  | x :: xs ->
```

```
    if Atomic.compare_and_set q old xs then Some x
```

```
    else pop_opt q
```



# Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t
```

```
let create (
```

```
let rec push
```

```
    let old =
```

```
    if Atomic.
```

```
    else push
```

```
let rec pop_
```

```
    let old =
```

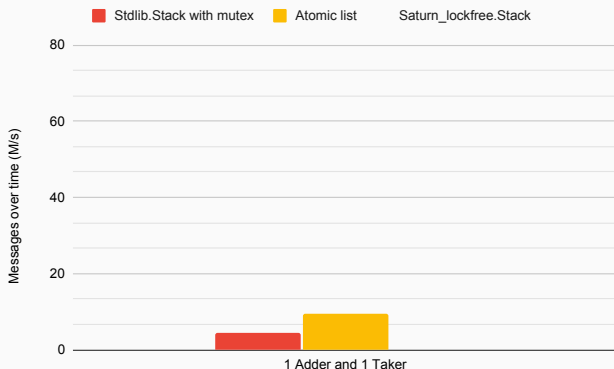
```
    match old
```

```
    | [] -> No
```

```
    | x :: xs ->
```

```
        if Atomic.compare_and_set q old xs then Some x
```

```
        else pop_opt q
```



# Concurrent stack: Treiber stack

```
type 'a t = 'a list Atomic.t
```

```
let create ()
```

```
let rec push
```

```
  let old = A
```

```
  if Atomic.c
```

```
  else push c
```

```
let rec pop_c
```

```
  let old = A
```


```
  match old v
```

```
  | [] -> Nor
```

```
  | x :: xs ->
```

```
    if Atomic.compare_and_set q old xs then Some x
```

```
    else pop_opt q
```



# Why Saturn?

- OCAML 5 : multicore programming
- Writting concurrent code is hard
  - ◇ bugs that can be hard to reproduce and understand
  - ◇ progress properties: deadlock, starvation etc..
- Writting efficient concurrent code is even harder
  - ◇ different behaviors on different CPU architectures
  - ◇ space to benchmarks is way larger
  - ◇ ...

## What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5

# What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
  - ◇ Queues
    - multi-producer, multi-consumer
    - single-producer, single-consumer
    - single-producer, multi-consumer
    - bounded, blocking (opened PRs)

# What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
  - ◇ Queues
    - multi-producer, multi-consumer
    - single-producer, single-consumer
    - single-producer, multi-consumer
    - bounded, blocking (opened PRs)
  - ◇ Work-stealing deque
  - ◇ Stacks
  - ◇ Hashtable (opened PR)
  - ◇ Skiplist: a sorted linked list with  $o(\log(n))$  operations
  - ◇ Bag



# What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
  - ◇ Queues
    - multi-producer, multi-consumer
    - single-producer, single-consumer
    - single-producer, multi-consumer
    - bounded, blocking (opened PRs)
  - ◇ Work-stealing deque
  - ◇ Stacks
  - ◇ Hashtable (opened PR)
  - ◇ Skiplist: a sorted linked list with  $o(\log(n))$  operations
  - ◇ Bag
- Most available data structures are lock-free

# What is in Saturn?

- A collection of concurrent-safe data structures for OCAML 5
  - ◇ Queues
    - multi-producer, multi-consumer
    - single-producer, single-consumer
    - single-producer, multi-consumer
    - bounded, blocking (opened PRs)
  - ◇ Work-stealing deque
  - ◇ Stacks
  - ◇ Hashtable (opened PR)
  - ◇ Skiplist: a sorted linked list with  $o(\log(n))$  operations
  - ◇ Bag
- Most available data structures are lock-free
  - Two libraries: SATURN and SATURN\_LOCKFREE

# Testing in Saturn

What should (and can) be tested ?

- Correctness
- Linearizability
- Progress (i.e. lock freedom)

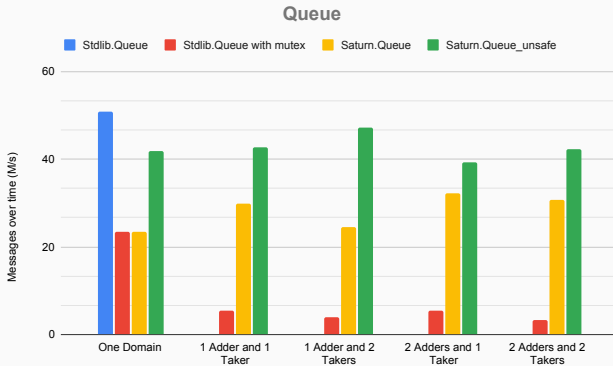
How ?

- Lin / STM (ocaml-multicore/multicoretest)
  - Correctness
  - Linearizability
- Dscheck (ocaml-multicore/dscheck)
  - Correctness
  - Lock-freedom


## Micro-optimizations

- ◇ On algorithms, to prevent false sharing, indirections etc..
- ◇ Experimental: some use **Obj.magic** (e.g. `Saturn.Queue_unsafe`)


# Is it any good ?



# Is it any good ?



# Is it any good ?



<https://github.com/lyrm/saturn-benchmarks/>

*Testing* a concurrent algorithm is hard due to the number of potential interleavings.

*Formal verification* is here to help us!





OCAML

---

Coq

$$\text{linearizability} \quad \underset{\text{in } \overline{\text{SC}}}{\approx} \quad \frac{\frac{\frac{\{ \text{stack-inv } t \}}{\langle \text{vs. stack-model } t \text{ vs } \rangle}}{\text{stack\_push } t \text{ } v}}{\langle \text{stack-model } t (v :: \text{vs}) \rangle}}{\{ () . \text{True} \}}$$

*Theorems for free from separation logic specifications*

Birkedal, Dinsdale-Young, Guéneau, Jaber, Svendsen & Tzevelekos

## Relaxed memory model (future work)

$$\frac{\frac{\frac{\{ \text{stack-inv } t \}}{\langle v_0, \dots, v_n, \mathcal{V}_0, \dots, \mathcal{V}_n. \text{stack-model } t ((v_0, \mathcal{V}_0), \dots (v_n, \mathcal{V}_n)) \rangle}}{\text{stack\_push } t \, v, \mathcal{V}}}{\frac{\langle \text{stack-model } t ((v, \mathcal{V}), (v_0, \mathcal{V}_0), \dots (v_n, \mathcal{V}_n)) \rangle}{\{ () . \text{True} \}}}$$

*Cosmo: a concurrent separation logic for multicore OCAML*

Mével, Jourdan & Pottier



Thank you for your attention!